


LoopCD at the same depth


Eight completed comparisons · gain labels are percentage points

Baseline Gain with LoopCD

Paper · Figure 1b



Our reproduction



AIME: 30 problems × 16 samples per arm. Our gain 95% problem-bootstrap intervals: 2.6B [4.79, 13.54]; 1.4B [0.83, 10.83].

HumanEval / MBPP: base-test pass@1. ARC-C / HellaSwag: length-normalized accuracy. Other uncertainty is in the result extracts.

* Looped-Qwen3 uses our independent reconstruction. Protocol differences remain; matching the direction does not mean matching the paper.

Source: arXiv:2610.02185v1 and results/figure1b.json · 2026-10-06